

# Long Latency Operations

D. Jiménez, E. Morancho and À. Ramírez

March 3, 2016

# Index

<b>Index</b>	<b>i</b>
<b>1 Previous Work</b>	<b>1</b>
1.1 Arithmetic Expression Optimizations . . . . .	1
1.2 Memoization . . . . .	2
<b>2 Arithmetic Expression Optimizations</b>	<b>1</b>
<b>3 Memoization and Buffering</b>	<b>2</b>
<b>4 Routine Specialization</b>	<b>3</b>

# 1

## Previous Work

We provide you a *.zip* file with several input files.

- There are several Makefile targets to automatically generate the binary files. For instance, *make prog.pg* will generate an executable compiled with profiling flags. Using *CFLAGS* variable in the Makefile in a subdirectory (local Makefile) of a program you can define the specific compilation flags for a program. Using *PROGS* variable in the local Makefile you can do "make" of different program names. You can read README file in the root directory of the laboratory lab3.
- Most of the programs have a parameter in order to modify the amount of work to do. You can modify that parameter in order to obtain a reasonable execution time.
- Most of the programs use a random input data. However, as they don't modify the seed of the random function, they should obtain the same results for different executions.

Some advices for the lab activities:

- You must redirect the output of the program executions to a file since the output can be large and/or with not printable characters.
- To compare ASCII output files, you should use tools like *diff*, *tkdiff* or *cmp*.
- To compare binary output files, you should use *cmp* tool.
- In order to visualize a file with non printable characters you can use the *od* tool (octal dump). For instance, `od -b /bin/ls | more` shows the ASCII code of the characters of the file */bin/ls*. That tool may be parametrized with the starting point to do the *dump*.

In order to take profit of the activities at the laboratory you should do the following activities before going to the lab:

### 1.1 Arithmetic Expression Optimizations

1. We give you the `primers.c` program that uses the Eratostenes method in order to create a prime list (from 1 upto the number you indicate as parameter).
  - Look at <http://www.wikipedia.org> for more information.
  - Try to find something regarding the `__udivdi3` and `__umoddi3`. That will help you with laboratory exercises. Hint: those symbols do not appear for 64-bit systems with a 64-bit OS and compiler.
    - Look for the GCC low-level runtime library.

## 1.2 Memoization

2. Look for information regarding to the trigonometric routine implementations `sin()` and `cos()`.
3. Look at the processor manual or the document we upload with the Lesson 3 in order to obtain the trigonometric instruction latencies.

## 2

# Arithmetic Expression Optimizations

1. `primers.c` program uses the Eratostenes method in order to create a prime list (from 1 upto the number you indicate as parameter).
  - (a) Compile `primers` with `"-O0"` (`make primers.o` and do timing, and keep the original output result).
  - (b) Compile with `make primers.pg` and use profiling to analyze its behavior when compiling with `-O0`.
  - (c) Which are the most time consuming functions? Look at their most time consuming lines to figure out which are the expensive operations. Is there any way to reduce/avoid those costly operations?
  - (d) Copy the original code to `primers_opt.c`. Optimize this copy of `primers` based on the analysis done in the previous question.
  - (e) Compile it using `make primers_opt.g` and check that the results of the optimized version you have done are the same than the original one.
  - (f) Which speedup have you obtained with your optimized code compared to the original `primers` when compiling both of them with `"-O0"`?
  - (g) Now, compile the original and the optimized versions of `primers` with `"-O3"` and do timing of both of them.
  - (h) Which speedup have you obtained compared to the original `primers` when compiling with `"-O3"` both programs, original and optimized? Do profiling again and look at the output of `objdump -d` in order to explain the speedup obtained.

# 3

## Memoization and Buffering

2. `trigon.c` program performs several calls to write and the trigonometric routines `sin()` and `cos()`.
  - (a) Study the source code of the program.
  - (b) Compile the program with `make trigon.pg3`. Profile and analyze the program behavior. Explain the results.
  - (c) Re-compile the program with `make trigon.pg3s`. Profile and analyze the program behavior. Explain the results.
  - (d) Figure out how many system calls are done in the program, and which is the elapsed time of the program.
  - (e) Based on the previous analysis, modify the program in order to avoid so many calls and reduce the overall elapsed time.
  - (f) Compute the speed-up of your optimized version compared to the original version.
  - (g) Profile the new code. How much CPU time is devoted to trigonometric computations? Which is the maximum speedup that we can achieve if we improve the corresponding execution time?
  - (h) Modify the program in order to avoid the computation repetitions and check that the output of the new version is the same as the original code.
  - (i) Profile the new version of the program and compare it with the profile of the original code.
  - (j) Which overall speed-up have you obtained?

## 4

# Routine Specialization

3. `pi.c` program compute the first 10000 decimals of the  $\pi$  number.

- (a) What is the speedup of the `pi` program compiled with `O3` compared to the `pi` compiled with `O0`?
- (b) Profile both `pi` binaries (compiled with `O3` and with `O0`) using `insmix_total` and `insmix_routines` scripts (`/assig/pcagrau/insmix_XXXX`) . And analyze the assembler code using `objdump -d`. Look for significant differences on the number of executed instructions for some type of instructions and open the assembler code to understand the differences. Could you justify these differences?
- (c) The compiler has done a good work but there is still work to do. Based on the previous profiling, look at the source code of the program, and in particular, the source code of routine `DIVIDE`. Then, propose an optimization that can help to reduce the cost of each `specific` call to `DIVIDE` in order to reduce/avoid long latency operations (Hint: Specilize the code of `pi` using memoization).
- (d) Profile the new version of the program and compare it with the profile of the original code.
- (e) Which speed-up have you obtained?